

python_stat

May 13, 2020

1 Statistical analysis

1.1 Introduction

Statistical analysis usually encompasses 3 activities in a data science workflow. These are (a) descriptive analysis, (b) hypothesis testing and (c) statistical modeling. Descriptive analysis refers to a description of the data, which includes computing summary statistics and drawing plots. Hypothesis testing usually refers to statistically seeing if two (or more) groups are different from each other based on some metrics. Modeling refers to fitting a curve to the data to describe the relationship patterns of different variables in a data set

In terms of Python packages that can address these three tasks:

Task	Packages
Descriptive statistics	pandas, numpy
Hypothesis testing	scipy, statsmodels
Modeling	statsmodels, lifelines

1.2 Descriptive statistics

Descriptive statistics that are often computed are the mean, median, standard deviation, inter-quartile range, pairwise correlations, and the like. Most of these functions are available in **numpy**, and hence are available in **pandas**. We have already seen how we can compute these statistics and have even computed grouped statistics. For example, we will compute these using the diamonds dataset

```
[2]: import numpy as np
import scipy as sc
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[3]: diamonds = pd.read_csv('data/diamonds.csv.gz')
```

```
[4]: diamonds.groupby('color')['price'].agg([np.mean, np.median, np.std])
```

```
[4]:
```

	mean	median	std
color			

D	3169.954096	1838.0	3356.590935
E	3076.752475	1739.0	3344.158685
F	3724.886397	2343.5	3784.992007
G	3999.135671	2242.0	4051.102846
H	4486.669196	3460.0	4215.944171
I	5091.874954	3730.0	4722.387604
J	5323.818020	4234.0	4438.187251

There were other examples we saw yesterday along these lines.

1.3 Hypothesis testing


Hypothesis testing is one of the areas where statistics is often used. There are functions for a lot of the standard statistical tests in `scipy` and `statsmodels`. However, I'm going to take a little detour to see if we can get some understanding of hypothesis tests using the powerful simulation capabilities of Python.

You have a coin and you flip it 100 times. You get 54 heads. How likely is it that you have a fair coin?

We can simulate this process, which is random, using Python. The process of heads and tails from coin tosses can be modeled as a **binomial** distribution. So we can repeat this experiment many many times on our computer, assuming we have a fair coin, and then see how likely what we observed is.

```
[9]: rng = np.random.RandomState(205)
x = rng.binomial(100, 0.5, 100000)
sns.distplot(x)
plt.axvline(54, color = 'r')
```

```
[9]: <matplotlib.lines.Line2D at 0x115c80580>
```



```
[10]: np.mean(x > 54)
```


```
[10]: 0.18456
```

This is what would be considered the *p-value* for the test that the coin is fair.

What happens if we increase the number of tosses, and we look at the proportion of heads. We observe 54% heads.

```
[11]: rng = np.random.RandomState(205)
x = rng.binomial(1000, 0.5, 100000)/1000
sns.distplot(x)
plt.axvline(0.54, color = 'r')
```

```
[11]: <matplotlib.lines.Line2D at 0x105b5ffd0>
```



```
[13]: np.mean(x > 0.54)
```

```
[13]: 0.00532
```

This is showing that as we increase the number of tosses, we get more robust signals about the coin. With 1000 tosses we see evidence that the coin is not fair, however small the unfairness is. This is related to the issue of statistical power. As we get more evidence, we can be more confident about what we are seeing.

We can use simulation to make better inference about differences between groups, or to get confidence intervals for a parameter of interest, without making assumptions about distributions or having to test for distributions. Python allows us to do that quickly and efficiently.

1.3.1 A permutation test

A permutation test is a 2-group test that asks whether two groups are different with respect to some metric. Let us look at a breast cancer proteomics experiment to illustrate this. The experimental data contains protein expression for over 12 thousand proteins, along with clinical data. We can ask, for example, whether a particular protein expression differs by ER status.

```
[14]: brca = pd.read_csv('data/brca.csv')
```

The idea about a permutation test is that, if there is truly no difference then it shouldn't make a difference if we shuffled the labels of ER status over the study individuals. That's literally what we will do. We will do this several times, and look at the average difference in expression each time. This will form the null distribution under our assumption of no differences by ER status. We'll then see where our observed data falls, and then be able to compute a p-value.


```
[23]: nsim = 10000
rng = np.random.RandomState(294)
x = np.where(brca['ER Status']=='Positive', -1, 1)
y = brca[protein].to_numpy()

obs_diff = np.nanmean(y[x==1]) - np.nanmean(y[x==-1])

diffs = np.zeros(nsim)
for i in range(nsim):
    x1 = rng.permutation(x)
    diffs[i] = np.nanmean(y[x1==1]) - np.nanmean(y[x1 == -1])
```

```
[24]: sns.distplot(diffs)
plt.axvline(x = obs_diff, color = 'r')
```

```
[24]: <matplotlib.lines.Line2D at 0x11b04d580>
```



```
[25]: pval = np.mean(np.abs(diffs) > np.abs(obs_diff))
f"P-value from permutation test is {pval}"
```

```
[25]: 'P-value from permutation test is 0.6709'
```

1.3.2 Testing many proteins

We could do the permutation test all the proteins using the array operations in numpy

```
[26]: expr_names = [u for u in list(brca.columns) if u.find('NP') > -1]

      exprs = brca[expr_names]
```

```
[32]: obs_diffs = exprs[x==1].mean(axis=0)-exprs[x==-1].mean(axis=0)
```

```
[34]: nsim = 1000
      diffs = np.zeros((nsim, exprs.shape[1]))
      for i in range(nsim):
          x1 = rng.permutation(x)
          diffs[i,:] =exprs[x1==1].mean(axis=0) - exprs[x1==-1].mean(axis=0)
```


```
[46]: pvals = np.zeros(exprs.shape[1])
      len(pvals)
```

```
[46]: 12395
```

```
[47]: for i in range(len(pvals)):
      pvals[i] = np.mean(diffs[:,i] > obs_diffs.to_numpy()[i])
```

```
[50]: sns.distplot(pvals)
```

```
[50]: <matplotlib.axes._subplots.AxesSubplot at 0x11a4866d0>
```



```
[51]: pd.Series(pvals).describe()
```

```
[51]: count      12395.000000
      mean         0.490345
      std          0.370710
      min          0.000000
      25%          0.105000
      50%          0.486000
      75%          0.868000
      max          1.000000
      dtype: float64
```

```
[54]: exprs_shortlist = [u for i, u in enumerate(list(exprs.columns)) if pvals[i] < 0.
      ↪0001 ]
```

```
[55]: len(exprs_shortlist)
```

```
[55]: 513
```

1.4 Regression analysis

The regression modeling frameworks in Python are mainly in `statsmodels`, though some of it can be found in `scikit-learn` which we will see tomorrow. We will use the diamonds dataset for demonstration purposes. We will attempt to model the diamond price against several of the other diamond characteristics

```
[58]: import statsmodels.api as sm
      import statsmodels.formula.api as smf

      mod1 = smf.glm('price ~ carat + clarity + depth + cut + color', data =
      ↪diamonds).fit()
```

```
[59]: mod1.summary()
```

```
[59]: <class 'statsmodels.iolib.summary.Summary'>
      """
              Generalized Linear Model Regression Results
      =====
      Dep. Variable:          price      No. Observations:          53940
      Model:                  GLM        Df Residuals:              53920
      Model Family:           Gaussian   Df Model:                  19
      Link Function:          identity   Scale:                     1.3382e+06
      Method:                  IRLS      Log-Likelihood:            -4.5699e+05
      Date:                   Wed, 13 May 2020      Deviance:                  7.2158e+10
      Time:                   13:21:00      Pearson chi2:              7.22e+10
      No. Iterations:          3
      Covariance Type:         nonrobust
      =====
      =====
```

	coef	std err	z	P> z	[0.025
0.975]					

Intercept	-6902.0434	245.309	-28.136	0.000	-7382.839
-6421.247					
clarity[T.IF]	5415.0087	52.191	103.754	0.000	5312.717
5517.301					
clarity[T.SI1]	3571.3831	44.613	80.053	0.000	3483.944
3658.822					
clarity[T.SI2]	2623.0139	44.813	58.532	0.000	2535.181
2710.846					
clarity[T.VS1]	4531.3874	45.570	99.437	0.000	4442.071
4620.704					
clarity[T.VS2]	4214.9672	44.865	93.948	0.000	4127.033
4302.901					
clarity[T.VVS1]	5068.3553	48.248	105.049	0.000	4973.792
5162.919					
clarity[T.VVS2]	4963.7218	46.924	105.781	0.000	4871.752
5055.692					
cut[T.Good]	644.1406	34.173	18.849	0.000	577.162
711.119					
cut[T.Ideal]	982.1534	31.780	30.905	0.000	919.866
1044.441					
cut[T.Premium]	849.9079	32.551	26.110	0.000	786.110
913.706					
cut[T.Very Good]	833.2870	32.291	25.806	0.000	769.998
896.576					
color[T.E]	-211.8364	18.316	-11.566	0.000	-247.734
-175.939					
color[T.F]	-303.2741	18.509	-16.385	0.000	-339.551
-266.997					
color[T.G]	-505.3602	18.127	-27.879	0.000	-540.888
-469.832					
color[T.H]	-977.5332	19.281	-50.699	0.000	-1015.323
-939.743					
color[T.I]	-1439.0796	21.655	-66.455	0.000	-1481.523
-1396.636					
color[T.J]	-2323.8709	26.731	-86.935	0.000	-2376.263
-2271.479					
carat	8885.8162	12.034	738.362	0.000	8862.229
8909.403					
depth	-7.1602	3.727	-1.921	0.055	-14.464
0.144					
=====					
=====					
"" ""					

This is the basic syntax for modeling in statsmodels. Let's go through and parse it.

One thing you notice is that we've written a formula inside the model

```
mod1 = smf.glm('price ~ carat + clarity + depth + cut + color', data = diamonds).fit()
```

This is based on another Python package, `patsy`, which allows us to write the model like this. This will read as “price depends on carat, clarity, depth, cut and color”. Underneath a lot is going on.

1. color, clarity, and cut are all categorical variables. They actually need to be expanded into dummy variables, so we will have one column for each category level, which is 1 when the diamond is of that category and 0 otherwise.
2. An intercept term is added
3. The dummy variables are concatenated to the continuous variables
4. The model is run

We can see the dummy variables using `pandas`

```
[60]: pd.get_dummies(diamonds)
```

```
[60]:
```

	carat	depth	table	price	x	y	z	cut_Fair	cut_Good	\
0	0.23	61.5	55.0	326	3.95	3.98	2.43	0	0	
1	0.21	59.8	61.0	326	3.89	3.84	2.31	0	0	
2	0.23	56.9	65.0	327	4.05	4.07	2.31	0	1	
3	0.29	62.4	58.0	334	4.20	4.23	2.63	0	0	
4	0.31	63.3	58.0	335	4.34	4.35	2.75	0	1	
...
53935	0.72	60.8	57.0	2757	5.75	5.76	3.50	0	0	
53936	0.72	63.1	55.0	2757	5.69	5.75	3.61	0	1	
53937	0.70	62.8	60.0	2757	5.66	5.68	3.56	0	0	
53938	0.86	61.0	58.0	2757	6.15	6.12	3.74	0	0	
53939	0.75	62.2	55.0	2757	5.83	5.87	3.64	0	0	

	cut_Ideal	...	color_I	color_J	clarity_I1	clarity_IF	clarity_SI1	\
0	1	...	0	0	0	0	0	
1	0	...	0	0	0	0	1	
2	0	...	0	0	0	0	0	
3	0	...	1	0	0	0	0	
4	0	...	0	1	0	0	0	
...
53935	1	...	0	0	0	0	1	
53936	0	...	0	0	0	0	1	
53937	0	...	0	0	0	0	1	
53938	0	...	0	0	0	0	0	
53939	1	...	0	0	0	0	0	

	clarity_SI2	clarity_VS1	clarity_VS2	clarity_VVS1	clarity_VVS2
0	1	0	0	0	0
1	0	0	0	0	0
2	0	1	0	0	0

3		0		0		1		0		0
4		1		0		0		0		0
...	
53935		0		0		0		0		0
53936		0		0		0		0		0
53937		0		0		0		0		0
53938		1		0		0		0		0
53939		1		0		0		0		0

[53940 rows x 27 columns]